

A DEEP LEARNING PROJECT

# FROM SCRATCH TO SMART

SERGIO EGUAKUN  
CHAITANYA DIWAKAR  
KATHERINE TORIAN



# OVERVIEW

GOALS &  
OBJECTIVES

DATASET  
SETUP

CNN MODEL

TRANSFER  
LEARNING

KEY  
TAKEAWAYS

# GOALS & OBJECTIVES

## BUILD A BASELINE CNN FOR IMAGE CLASSIFICATION

Establish a simple benchmark model to understand CIFAR-10 and measure improvements against it.

## EXPLORE STRATEGIES TO IMPROVE CNN PERFORMANCE

Apply techniques like data augmentation, batch normalization, dropout, and L2 regularization to enhance feature learning and reduce overfitting.

## LEVERAGE TRANSFER LEARNING WITH PRETRAINED MODELS

Adapt state-of-the-art architectures (MobileNetV2, EfficientNetB0) to CIFAR-10 and compare their effectiveness against custom CNNs.

## EVALUATE AND COMPARE MODELS TO IDENTIFY THE BEST APPROACH

Systematically assess test accuracy across models to highlight the impact of architectural choices and demonstrate how transfer learning achieves top performance (EfficientNetB0: ~93%).



# Setup Setup Setup Setup Dataset



## OUR DATASET

- 60,000 color images ( $32 \times 32$ , RGB) across 10 classes (airplane, car, bird, cat, deer, dog, frog, horse, ship, truck).
- 50,000 training images + 10,000 test images.

## DATA SPLITS & LABELS

- Used `cifar10.load_data()` for train/test split.
- CNN models: integer labels + `sparse_categorical_crossentropy`.
- Transfer learning models: one-hot labels + `categorical_crossentropy`.

## PREPROCESSING

Pixel values scaled to [0,1] (default) or preprocessed via model-specific functions (e.g., `preprocess_input` for MobileNetV2, EfficientNet).

## EVALUATION METRIC:

Accuracy (training, validation, and test).

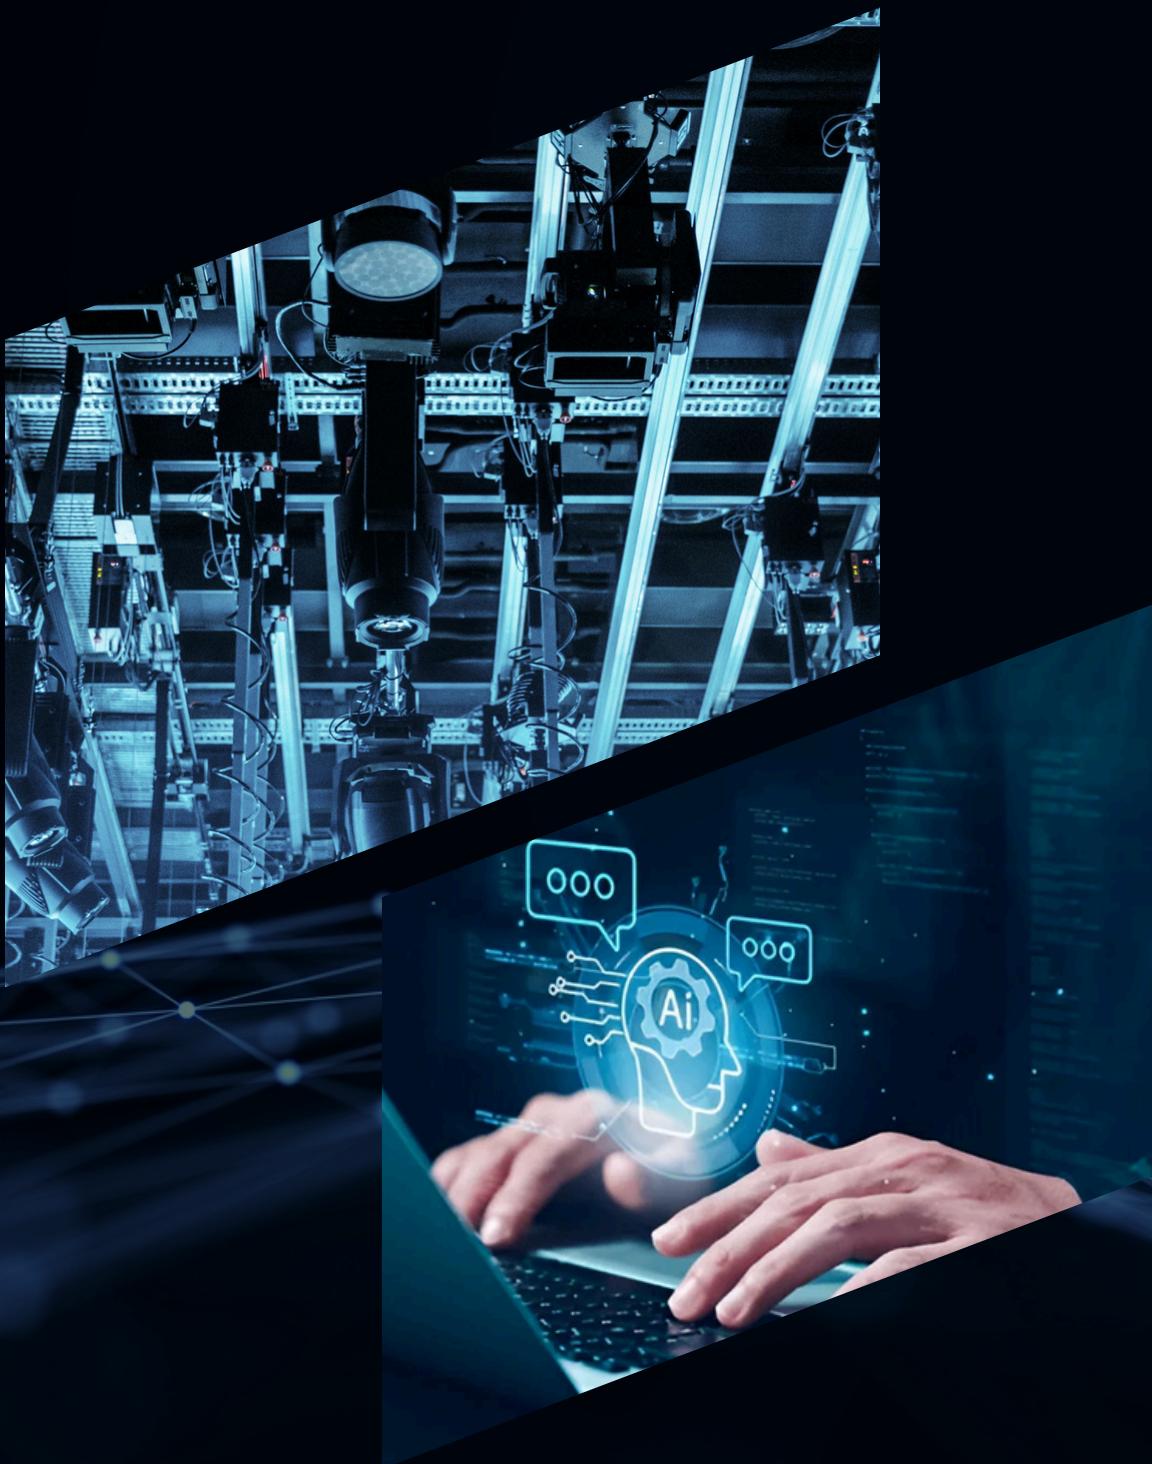
# OUR BASELINE CNN MODEL

		Predicted									
		airplane	automobile	bird	cat	deer	dog	frog	horse	ship	truck
True	airplane	672	20	59	24	23	1	17	19	141	24
	automobile	30	757	13	10	7	16	13	8	66	80
	bird	80	9	569	84	82	33	60	52	25	6
	cat	41	9	89	481	69	137	88	49	19	18
	deer	24	5	92	80	620	34	55	67	18	5
	dog	25	3	83	218	53	476	41	83	7	11
	frog	10	6	64	65	56	25	751	6	12	5
	horse	23	6	53	46	70	44	15	731	5	7
	ship	53	31	10	23	11	4	4	7	836	21
	truck	37	101	14	24	5	11	13	35	71	689

WE CAN CLEARLY SEE THE MISPREDICTIONS OF THE BASELINE MODEL.

THAT'S WHY WE HAVE TO REFINER OUR MODEL AND MAKE IT AS ACCURATE AS POSSIBLE.

# CNN MODELS



## 1- BASELINE CNN

- Architecture: Conv32 → MaxPool → Conv64 → MaxPool → Flatten → Dense128 → Dense10
- Features: Simple, interpretable, uses sparse categorical crossentropy
- Test Accuracy: 67.7% | Training Accuracy: 98.5% | Validation Accuracy: 75.1%
- Takeaway: Overfits easily, limited feature extraction

## 2- BASELINE CNN + EARLY STOPPING + DATA AUGMENTATION

- Techniques: EarlyStopping (patience=3), Data Augmentation (flip, shift, zoom)
- Test Accuracy: 69.5% (ES), 74.2% (ES + Augmentation) Training Accuracy: 75.6% | Validation Accuracy: 78.1%
- Takeaway: Augmentation significantly improves generalization

## 3- STRONGER CUSTOM CNN (MODEL 2)

- Architecture: 3 blocks: Conv → BatchNorm → Conv → BatchNorm → MaxPool → Dropout
- Filters: 32 → 64 → 128, Dense256 before classifier
- Regularization: L2 + Dropout
- Test Accuracy: 84.3% | Training Accuracy: 89.5% | Validation Accuracy: 84.1%
- Takeaway: Deeper network + BatchNorm + regularization reduces overfitting and captures richer features



# TRANSFER LEARNING

## MOBILENETV2

- Input resized:  $32 \times 32 \rightarrow 96 \times 96$
- Pretrained on ImageNet, base frozen
- Custom head: GAP → Dropout(0.3) → Dense10
- Test Accuracy: 85.6% | Training Accuracy: 84.8% | Validation Accuracy: 85.8%
- Key point: Lightweight model, pretrained features generalize well

## EFFICIENTNETB0

- Input resized:  $32 \times 32 \rightarrow 224 \times 224$
- Labels one-hot, base pretrained on ImageNet
- Training:
  - Phase 1: freeze base, train head (10 epochs, LR=1e-4)
  - Phase 2: fine-tune last ~30 layers (10 epochs, LR=1e-5)
- Custom head: GAP → Dropout0.3 → Dense10
- Test Accuracy: 92.7% | Training Accuracy: 91.8% | Validation Accuracy: 92.6%
- Key point: Fine-tuning pretrained features gives best generalization



# OUR FINAL REFINED MODEL

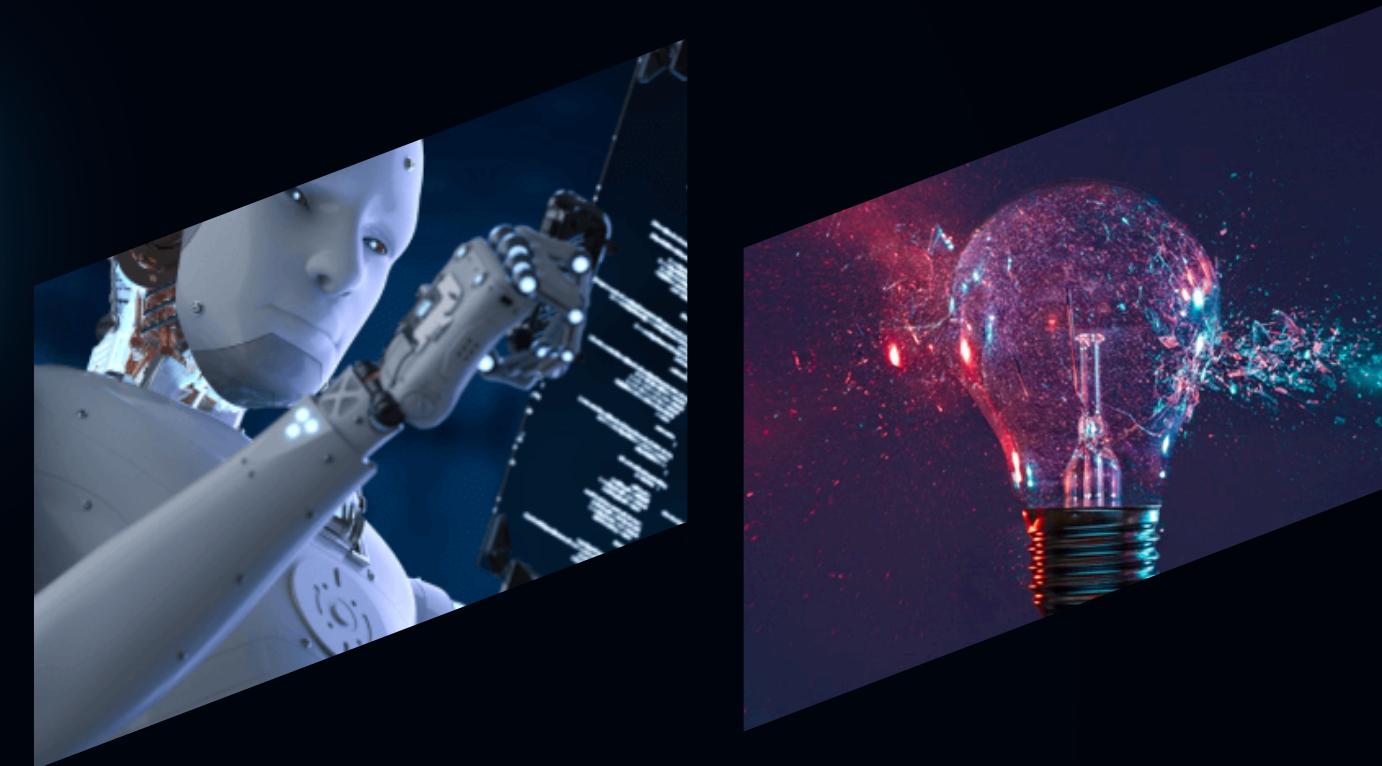
		Predicted									
		automobile	bird	cat	deer	dog	frog	horse	ship	truck	
True	airplane	785	11	46	19	16	4	12	17	45	45
	automobile	10	883	1	2	2	4	10	2	9	77
	bird	69	3	579	49	105	72	89	24	3	7
	cat	18	5	47	565	60	194	63	32	4	12
	deer	15	2	29	53	761	32	52	50	4	2
	dog	7	2	26	125	29	747	18	38	2	6
	frog	4	0	18	43	25	18	879	6	3	4
	horse	8	2	21	34	46	54	6	817	1	11
	ship	57	31	9	15	3	4	5	6	847	23
	truck	17	46	4	14	2	4	7	13	10	883

WITH A SUPER HIGH ACCURACY, WE CAN NOW SEE HOW OUR MODEL IMPROVED AND BECAME MORE PRECISE AFTER ALL THE REFINING.

# KEY TAKEAWAYS

## EARLY STOPPING

Prevents overfitting, small improvement



## DATA AUGMENTATION

Improves generalization significantly (+5% accuracy).

## BATCHNORM + L2 + DROPOUT

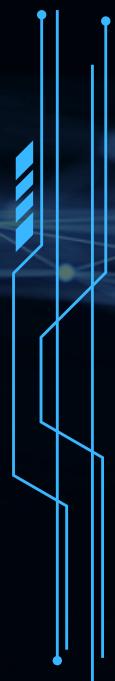
Deeper, regularized CNN reduces overfitting and captures richer features (+10% accuracy).

## MOBILENETV2

Pretrained features outperform custom CNNs (+0.6 points vs Model 2), lightweight and efficient.

## EFFICIENTNETB0

Fine-tuned transfer learning achieves state-of-the-art accuracy (92–93%), best generalization.



# ABOUT THE FUTURE...

- Experiment with larger or more advanced architectures (e.g., EfficientNetB3/B4) for potential accuracy gains.
  - Explore semi-supervised or self-supervised learning to reduce dependency on labeled data.
- Apply transfer learning strategies to other small or domain-specific datasets.
  - Combine multiple models (ensembles) to push accuracy even higher.
- Investigate optimization techniques and learning rate schedules for more efficient training.



# THANK YOU

ANY QUESTIONS?

